Parallel Computing - Part 2 - GPU

R. Guivarc'h

Ronan.Guivarch@toulouse-inp.fr

2025 - ModIA

Acknowledgements

Used materials and examples from

- **Introduction to GPU Programming with the OpenMP API, Dr.-Ing. Michael Klemm, OpenMP Architecture Review Board**
- Introduction aux GPGPU, Amina Guermouche, cours de Télécom SudParis
- OpenMP on GPUs, First Experiences and Best Practices, Jeff Larkin, NVIDIA
- OpenMP 5.0/5.1 Tutorial, Exascale Computing Project
- Introduction to OpenACC and OpenMP GPU, Pierre-François Lavallée, Thibaut Véry, IDRIS
- Introduction to Directive Based Programming on GPU, Helen He, NERSC
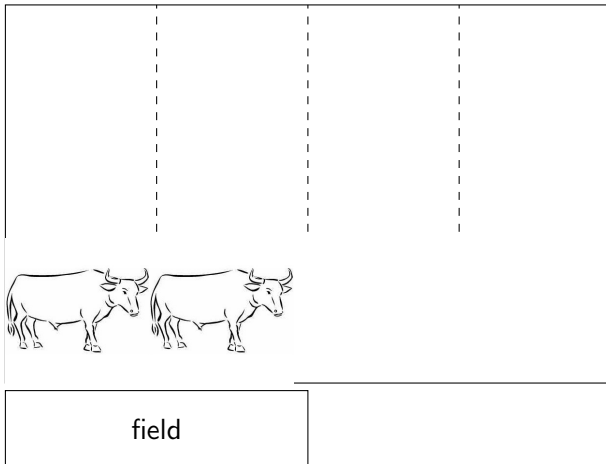- Applications of GPU Computing, Alex Karantza

Web Pointers

- **Intro to GPU Programming with the OpenMP API (OpenMP Webinar)**
- Introduction to OpenMP - Tim Mattson (Intel)
- GPU, IA et Big Data : comprendre le rôle des cartes graphiques en Data Science
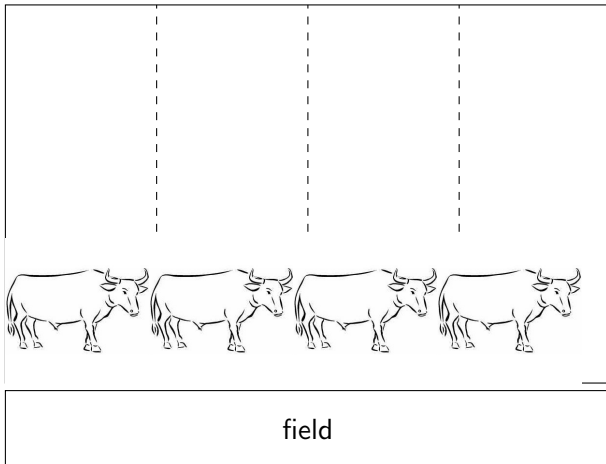- GPUs explained
- NVIDIA, AMD

A bit of history
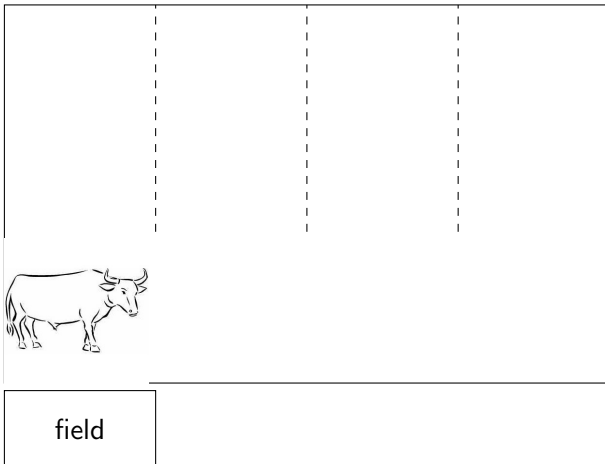
field
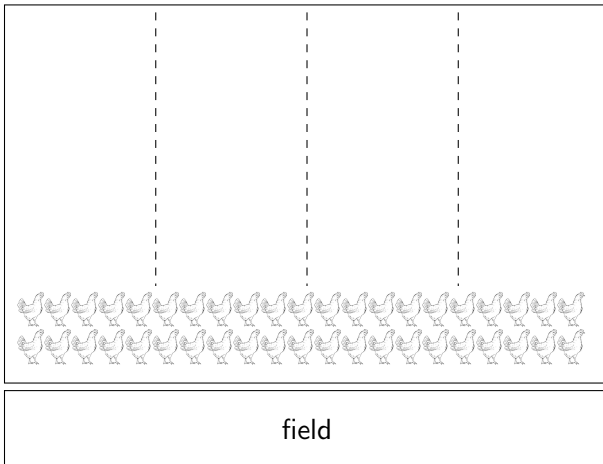
A bit of history

- Increased number of CPUs
- ☺ Faster

A bit of history

- Increased number of CPUs
- 🙂 Faster
- 🙁 More expensive, physical limits
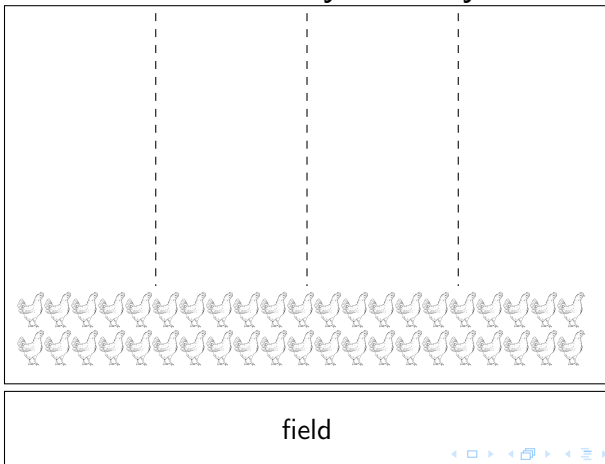
a bit of history

field

a bit of history

"If you were plowing a field, which would you rather use: Two strong oxen or 1024 chickens?" **Seymour Cray**

field

CPU VS GPU: demo

https://www.youtube.com/watch?v=-P28LKWTzrI

GPUs in Industry

Many applications have been developed to use GPUs for supercomputing in various fields

- Scientific Computing: CFD, Molecular Dynamics, Genome Sequencing, Mechanical Simulation, Quantum Electrodynamics
- Image Processing: Registration, interpolation, feature detection, recognition, filtering
- Data Analysis: Databases, sorting and searching, data mining

Major Categories of Algorithm

- 2D/3D filtering operations
- n-body simulations
- Parallel tree operations – searching/sorting
- ...

All suited to GPUs because of data-parallel requirements and **uniform** kernels

Computational Fluid Dynamics

- Simulate fluids in a discrete volume over time
- Involves solving the Navier-Stokes partial differential equations iteratively on a grid (can be considered as a filtering operation)
- When parallelized on a GPU using multigrid solvers, 10x speedups have been reported
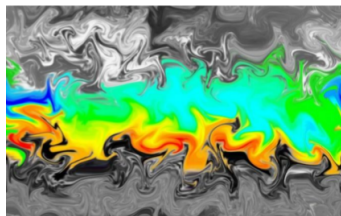
Molecular Dynamics

- Large set of particles with forces between them – protein behavior, material simulation
- Calculating forces between particles can be done in parallel for each particle
- Accumulation of forces can be implemented as multilevel parallel sums
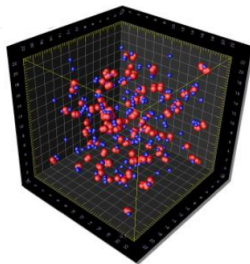
Genetics

- Large strings of genome sequences must be searched through to organize and identify samples
- GPUs enable multiple parallel queries to the database to perform string matching
- Again, order of magnitude speedups reported
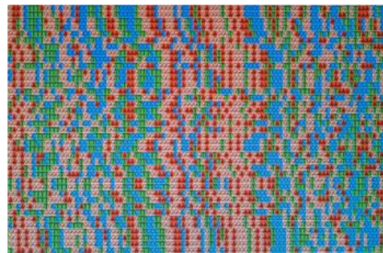
Electrodynamics

- Simulation of electric fields, Coulomb forces
- Requires iterative solving of partial differential equations
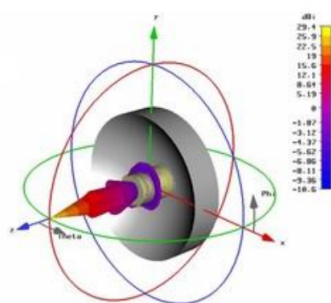- Cell phone modeling applications have reported 50x speedups using GPUs

Image Processing

- Medical Imaging was the early adopter
  - Registration of massive 3D voxel images
  - Both the cost function for deformable registration and interpolation of results are filtering operations
- Generic feature detection, recognition, object extraction are all filters
- For object recognition, one can search a database of objects in parallel
- Running these algorithms off the CPU can allow real-time interaction
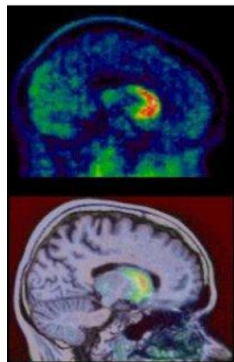
Data Analysis

- Huge databases for web services require instant results for many simultaneous users
- Insufficient room in main memory, disk is too slow and doesn't allow parallel reads
- GPUs can split up the data and perform fast searches, keeping their section in memory

Artificial Intelligence and Big Data

- Some GPUs are efficient for matrix multiplication and convolution (tensor cores)
- Bandwidth-optimised, ability to process a large amount of data
- Machine Learning models can be trained 215 times faster.

CPU vs. GPU

- Latency vs. Throughput
- Task parallelism vs. Data parallelism
- Multi-threaded vs. SIMD
- Tens of threads vs. Tens of thousands of threads

Latency & Throughput

- Latency is the delay between the time an operation is initiated, and the moment when its effects become detectable
  - A car has a lower latency than a bus (faster)
- Throughput is the amount of work done over time
  - A bus has a higher throughput than a car (more people at a time)

Latency & Throughput

- CPUs must minimize latency (neglecting throughput ☹)
  - keyboard input
- CPUs maximize out-of-cache operations (pre-fetch, out-of-order execution, …)
- → CPUs need large caches

Latency & Throughput

- CPUs must minimize latency (neglecting throughput ☹)
    - keyboard input
- CPUs maximize out-of-cache operations (pre-fetch, out-of-order execution, …)
- → CPUs need large caches
- GPUs are high performance, high throughput processors
    - They do not need large caches
    - More transistors can be dedicated to computation

Chipset

- More transistors are dedicated to data processing instead of cache management
- Chip of the same size but with more ALUs, so more threads for computation
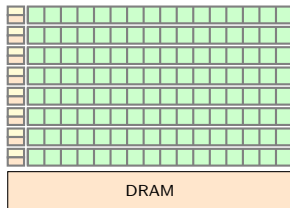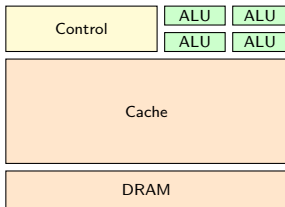
Thread management on GPUs

- **How manage**
  - Synchronisation between as many threads
  - Scheduling, context switching
 - Programmation

Thread management on GPUs

- **How manage**
  - Synchronisation between as many threads
  - Scheduling, context switching
  - Programmation
- threads on GPUs are:
  - Independent (no synchronisation)
  - SIMD (reduced scheduling cost)
  - Programming by block of threads

CPU Parallelism vs. GPU Parallelism

CPU: Task Parallelism

- Simultaneous execution of several functions on different cores and on different data, or not
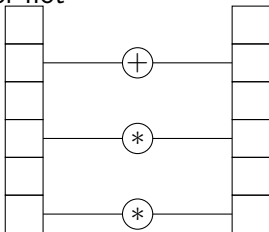
GPU: Data Parallelism

- Simultaneous execution of the same function by several cores on different data
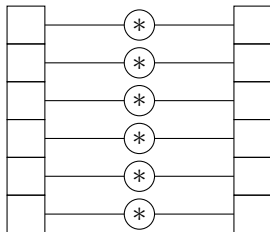
Stream processing

- The fundamental unit of a GPU is the stream processor (SP)
  - Large amount of data (*"stream"*)
  - Execute the same operation ("kernel" or "shader") on all the data
- we have also SM (stream multi-processor)

Architecture of a converged node on Jean Zay

CNRS/IDRIS-GENCI

Architecture of a converged node on Jean Zay

The number of compute cores on machines with GPUs is much greater than on a classical machine.

IDRIS' Jean-Zay has 2 partitions:

- Non-accelerated : $2\times20 = 40$ cores   1528 nodes (CPU)
- Accelerated with 4 Nvidia V100 = $32\times80\times4$ = 10240 cores 612 nodes (CPU+GPU)
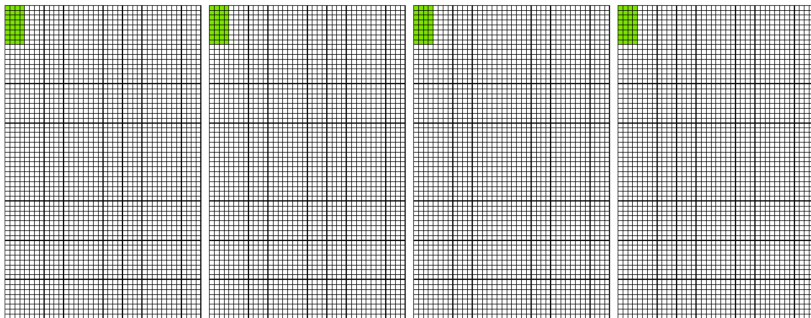
Architecture of a converged node on Jean Zay

Architecture of a converged node on Jean Zay (40 cores, 192 GB of memory, 4 GPU V100)
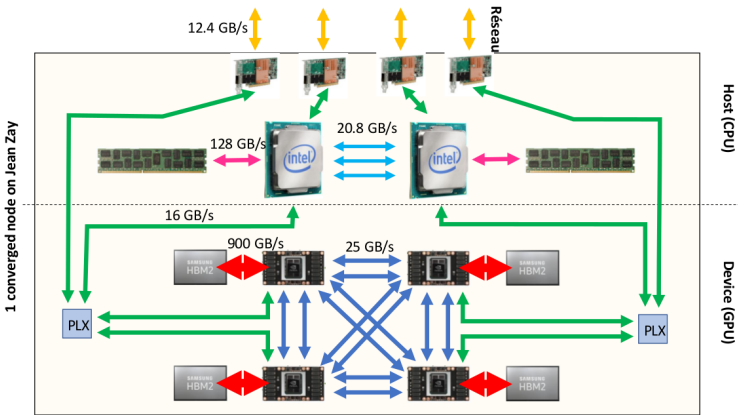
General Architecture of a converged node

Programming GPUs

Programming effort and technical expertise ⟩

Libraries
- cuBLAS
- cuSPARSE
- AmgX
- MAGMA
- pytorch

Directives
- OpenACC
- OpenMP 6.0 (nov. 2024)

Programming languages
- CUDA
- OpenCL
- HIP

Programming GPUs

Programming effort and technical expertise ⟶

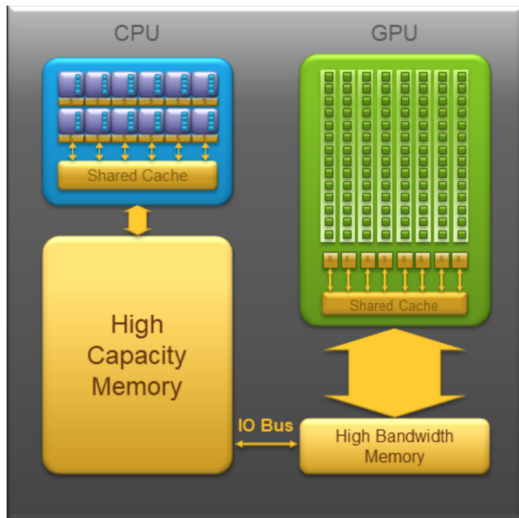| Libraries	Directives	Programming languages
○ cuBLAS	○ OpenACC	○ CUDA
○ cuSPARSE	○ OpenMP 6.0 (nov. 2024)	○ OpenCL
○ AmgX		○ HIP
○ MAGMA		
○ pytorch		

+ Minimum change in the code	+ Easy to use	- Harder to use
+ Free performance	+ Portable	- Need to rewrite application
- Limited by the available libraries and their functionalities	- Performance	- Less portable
	(user has less control)	++ Fine control of performance & flexibility

Directives

OpenMP target

- http:
  //www.openmp.org/
- First standard OpenMP
  4.5 (11/2015)
- Latest standard OpenMP
  6.0 (11/2024)
- Main compilers:
  - Cray (for Cray
    hardware)
  - GCC (since 7)
  - CLANG
  - IBM XL
  - PGI

OpenAcc

- http:
  //www.openacc.org/
- Cray, NVidia, PGI, CAPS
- First standard 1.0
  (11/2011)
- Latest standard 3.0
  (11/2019)
- Main compilers:
  - PGI
  - Cray (for Cray
    hardware)
  - GCC (since 5.7)

Sample OpenMP and OpenACC Codes

```
#define N 128
double x[N*N];
int i, j, k;
for (k=0; k<N*N; ++k) x[k] = k;

#pragma omp target
#pragma omp teams distribute
for (i=0; i<N; ++i) {
#pragma omp parallel for simd
  for (j=0; j<N; ++j) {
    x[j+N*i] *= 2.0;
  }
}
```

```
#define N 128
double x[N*N];
int i, j, k;
for (k=0; k<N*N; ++k) x[k] = k;

#pragma acc parallel
#pragma acc gang worker
for (i=0; i<N; ++i) {
#pragma acc vector
  for (j=0; j<N; ++j) {
    x[j+N*i] *= 2.0;
  }
}
```

directives

CUDA code example (for information)

kernel

```
__global__ void add (int *a, int *b, int *c){
  c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];
}
```

main

```
#define N 512
int main (void){
  int *a, *b, *c;
  int *gpu_a, *gpu_b, *gpu_c;
  int size = N * sizeof(int);
  // memory allocation on GPU                 // allocation on CPU
  cudaMalloc((void **)&gpu_a, size);          a=(int*) malloc (size);
  cudaMalloc((void **)&gpu_b, size);          b=(int*) malloc (size);
  cudaMalloc((void **)&gpu_c, size);          c=(int*) malloc (size);
```

CUDA code example (for information)

main (cont.)

```
// Initialization of a and b
...
// Data transfert to GPU
cudaMemcpy (gpu_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy (gpu_b, b, size, cudaMemcpyHostToDevice);

// call the kernel on N threads
add <<< N, 1 >>> (gpu_a, gpu_b, gpu_c);

// Data transfert from GPU
cudaMemcpy(c, gpu_c, size, cudaMemcpyDeviceToHost);

// Free memory (both CPU and GPU)
free(a); free(b); free(c);
cudaFree(gpu_a);
cudaFree(gpu_b);
cudaFree(gpu_c);
return 0;
}
```

Advantages of Directive Based Parallelism

- Incremental parallel programming
  - Find hotspot, parallelize, check correctness, repeat
- Single source code for sequential and parallel programs
  - Use compiler flag to enable or disable
  - No major overwrite of the serial code
- Works for both CPU and GPU
- Low learning curve, familiar C/C++/Fortran program environment
  - Do not need to worry about lower level hardware details
- Simple programming model than lower level programming models
- Portable implementation:
  - Different architectures, different compilers handle the hardware differences

Host $\Leftrightarrow$ Device Model
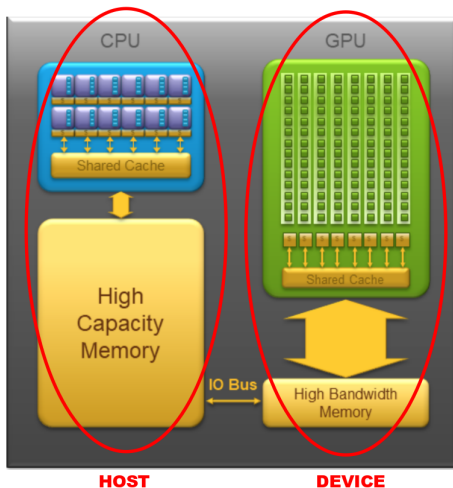
Device Execution Model

- Device: an implementation-defined logical execution unit
- Can have a single host and one or more target devices (accelerators)
- Host and Device have separate data environment (except with managed memory or unified shared memory)
- The execution model is host-centric
    - Host creates data environment on the device(s)
    - Host maps data to the device data environment.
    - Host then offloads accelerator regions to the device for execution
    - Host updates the data between the host and the device.
    - Host destroys data environment on device.

CUDA Interoperability

- OpenMP is a high-level language, sometimes low level optimizations will be necessary for best performance
- CUDA Kernels or Accelerated libraries for example
- The use_device_ptr map type allows OpenMP device arrays to be passed to CUDA or accelerated libraries
- The is_device_ptr map clause allows CUDA arrays to be used within OpenMP target regions

Example of `use_device_ptr`

```
#pragma omp target data map(alloc:x[0:n]) map(from:y[0:n]) ⟵
{
    #pragma omp target teams distribute parallel for
    for( i = 0; i < n; i++)
    {
        x[i] = 1.0f;
        y[i] = 0.0f;
    }

    #pragma omp target data use_device_ptr(x,y)
    {
        cublasSaxpy(n, 2.0, x, 1, y, 1);
    }
}
```

Manage data
movement using
map clauses

Expose the device
arrays to CUBLAS

GPU Programming with the OpenMP API

OpenMP Webinar

> Intro to GPU Programming with the OpenMP API
> https://www.openmp.org/events/
> intro-to-gpu-programming-with-the-openmp-api/